

BANZAICLOUD

Reconciliation on different levels

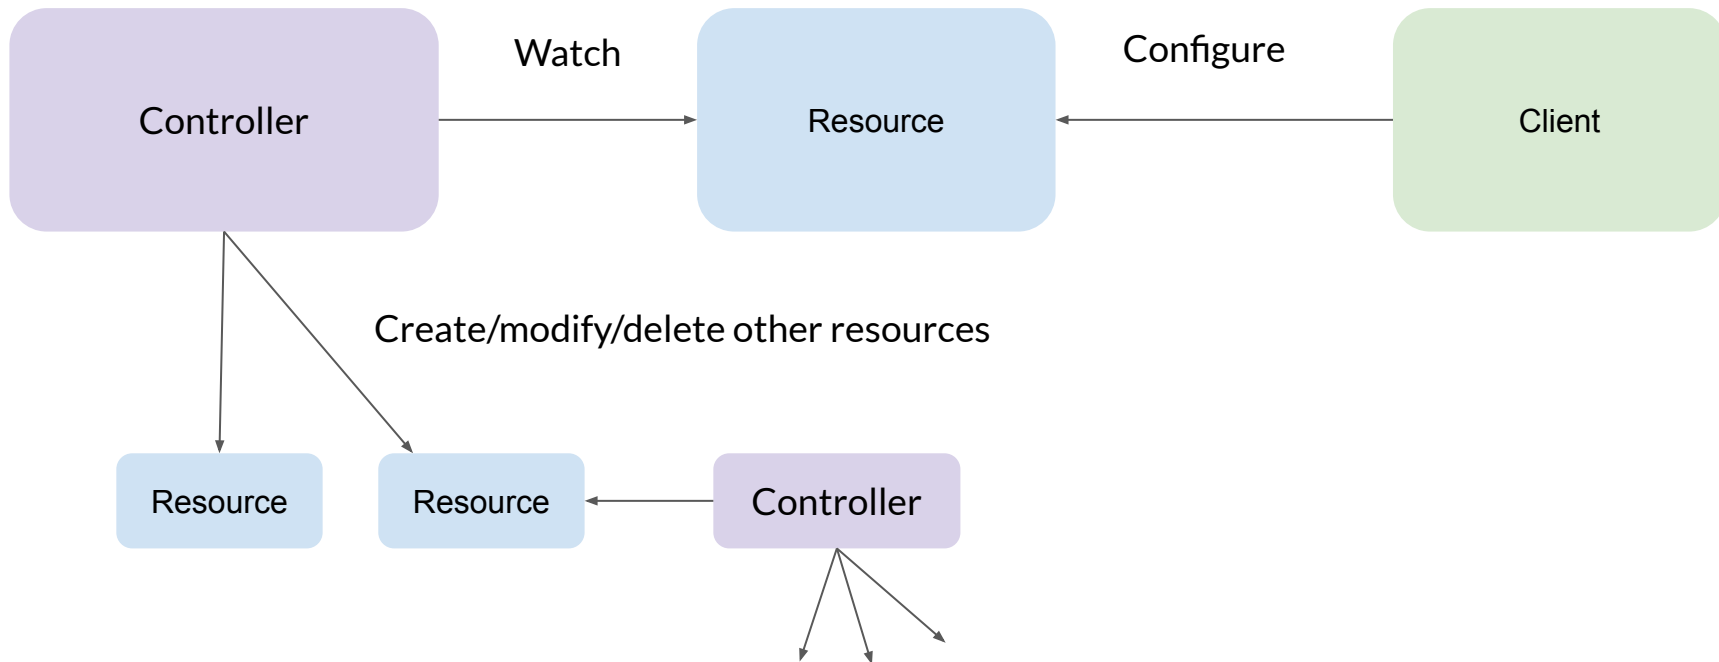
Peter Wilcsinszky (pepov)

banzaicloud.com

Reconcile



Reconcile





<https://github.com/banzaicloud/thanos-operator>

apiVersion: monitoring.banzaicloud.io/v1alpha1


kind: Thanos

metadata:


name: exampe

spec:

query: {}



Deterministic names instead of storing IDs



<https://github.com/banzaicloud/operator-tools/tree/master/pkg/reconciler>

type ResourceBuilder **func**(runtime.Object) (runtime.Object, DesiredState)

type ResourceReconciler **interface** {
 ReconcileResource(runtime.Object, DesiredState) (*reconcile.Result, error)
}



```
type StaticDesiredState string
```

```
const (  
    StateAbsent StaticDesiredState = "Absent"  
    StatePresent StaticDesiredState = "Present"  
)
```

```
type DesiredState interface {  
    BeforeUpdate(current runtime.Object) error  
}
```


How do we know we have to update a resource?

Can't we update it all the time?



Kubernetes API server would handle these updates, but...

There would be lots of unnecessary connections to the API server

We lose the visibility into what we are doing exactly, i.e. the ability to track down inconsistencies

Get the existing resource (from the cache)

Compare with the original version using


<https://github.com/banzaicloud/k8s-objectmatcher>

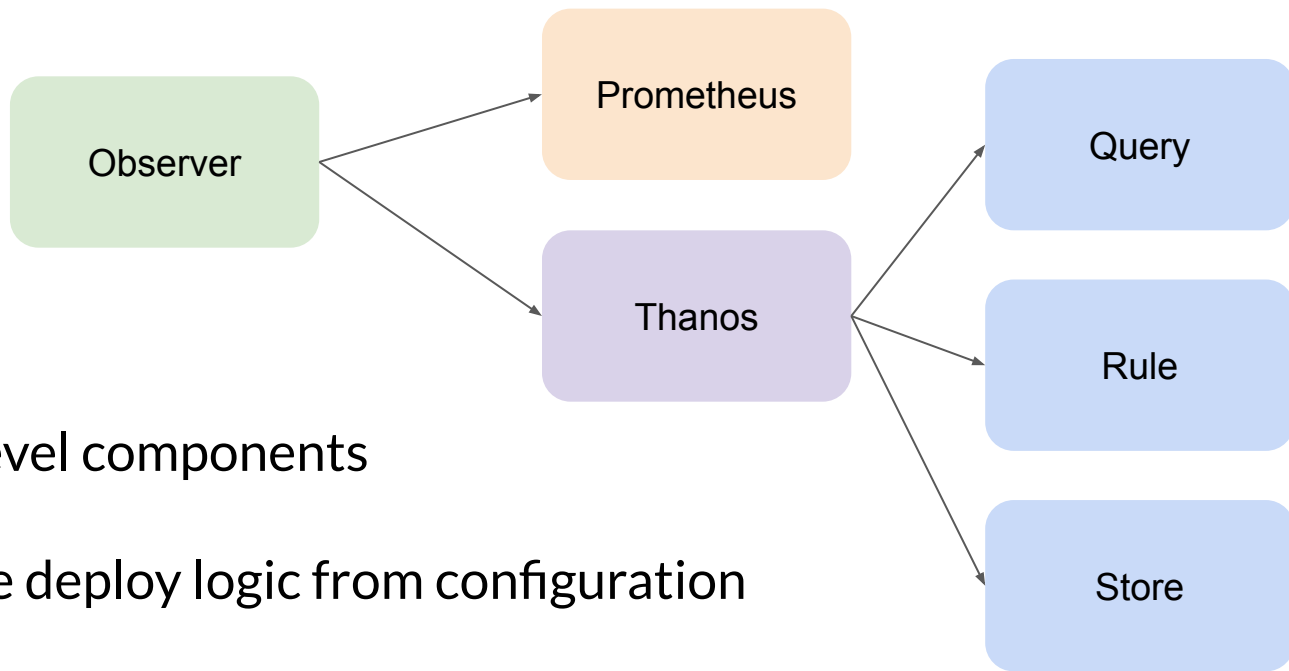
Update only if there is a diff



Component

A bunch of resources managed together





Higher level components

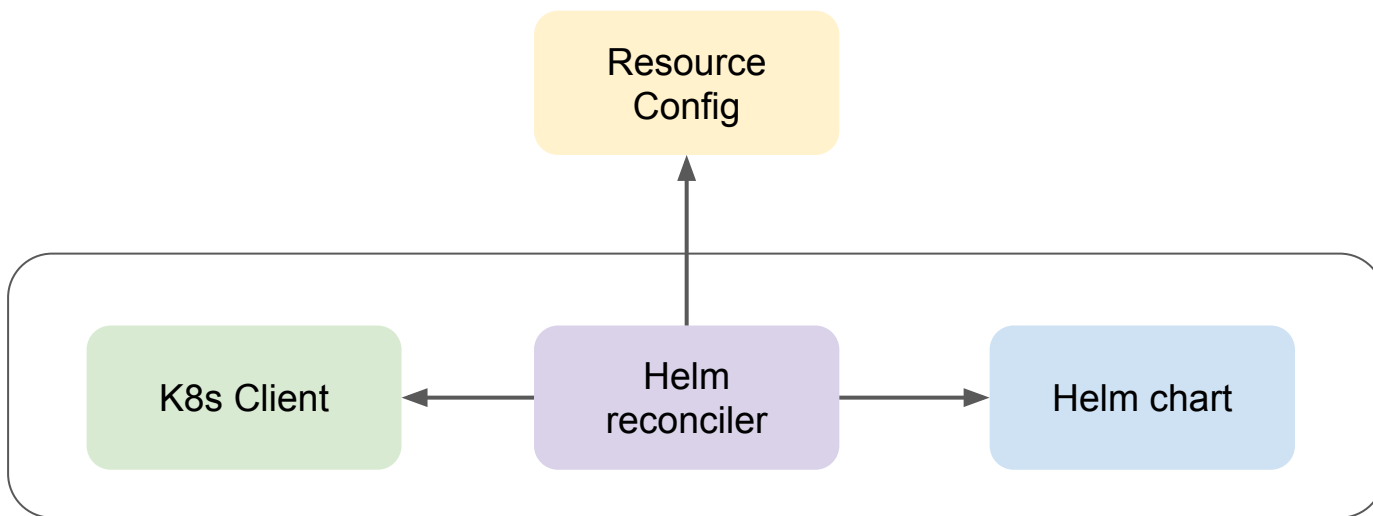
Decouple deploy logic from configuration

Possible implementations

- Helm v3
- External resource builder
- `whatnot | kubectl apply -f-`

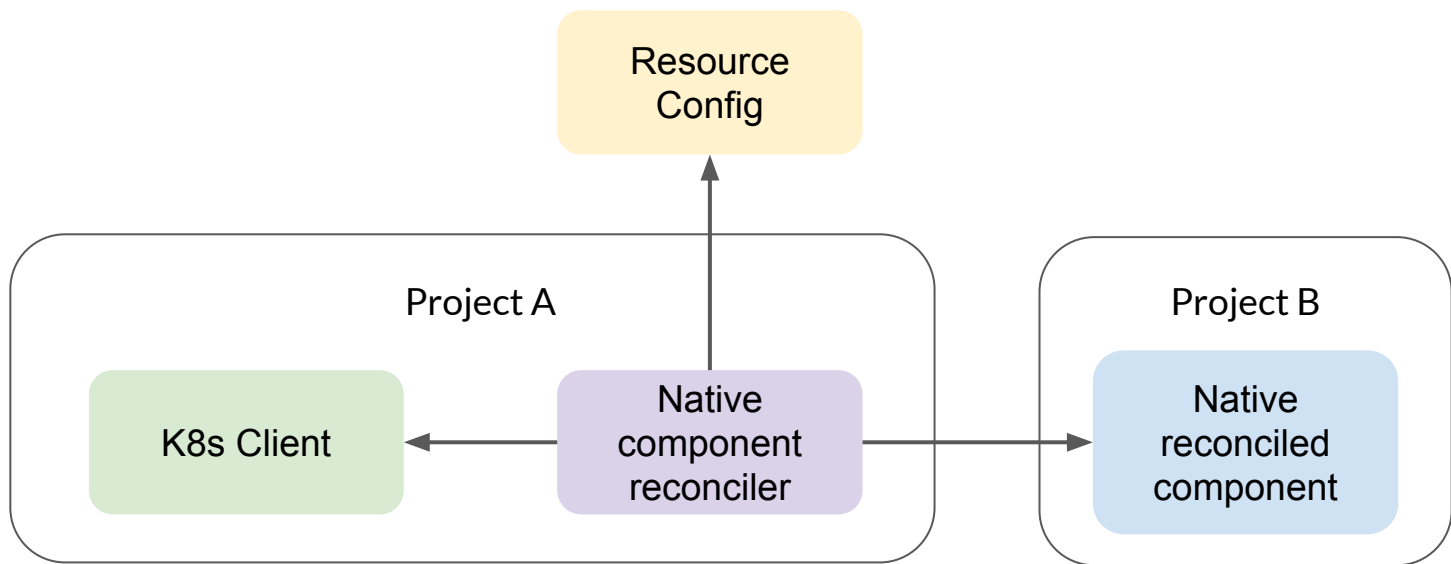
Pro: use existing Helm charts

Con: little to no control over the configuration



Pro: tight control over resources, flexible object manipulation

Con: need to write code for it





```
func ResourceBuilders(parent reconciler.ResourceOwner, object interface{}) []reconciler.ResourceBuilder {
    config := &ComponentConfig{
        Disabled: true,
    }
    if object != nil {
        config = object.(*ComponentConfig)
    }
    resources := []reconciler.ResourceBuilder{
        config.build(parent, Operator),
        config.build(parent, ClusterRole),
        config.build(parent, ClusterRoleBinding),
        config.build(parent, ServiceAccount),
    }
    // We don't return with an absent state since we don't want them to be removed
    if !config.Disabled {
        resources = append(resources,
            func() (runtime.Object, reconciler.DesiredState, error) {
                return CRD(config, v1alpha1.GroupVersion.Group, "objectstores")
            },
            func() (runtime.Object, reconciler.DesiredState, error) {
                return CRD(config, v1alpha1.GroupVersion.Group, "thanos")
            },
            func() (runtime.Object, reconciler.DesiredState, error) {
                return CRD(config, v1alpha1.GroupVersion.Group, "storeendpoints")
            },
        )
    }
    return resources
}
```

<https://banzaicloud.com/products/one-eye/>

apiVersion: one-eye.banzaicloud.io/v1alpha1

kind: Observer

metadata:

name: thanos


spec:


thanos:

operator: {}

prometheus:

prometheusOperatorChart: {}





BANZAI **CLOUD**